

Predicting Loan Defaults Based on Customer Behavior

Case Study

- This provides a case study of "a "Predicting Loan Defaults Based on Customer Behavior" project.
- by Li-Hsuan, Lin

1. The Data Analysis Problem

- Loan Defaults prediction data set
- The purpose is to predict whether the company should approve an application for loan by a given individual.

Variable	Description
Id	Unique identifier
Income	Income
Age	Age
Experience	Years of work experience
Married/Single	Marital status (Married/Single)
House_Ownership	House ownership status
Car_Ownership	Car ownership status
Profession	Profession or occupation
CITY	City
STATE	State
CURRENT_JOB_YRS	Years in current job
CURRENT_HOUSE_YRS	Years in current residence
Risk_Flag	Risk flag indicating potential risk (1 for risk, 0 for no risk)

2. Exploring the Data Using Pandas

2.1 Import the necessary modules and read the data

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns # a visualisation library we have
from matplotlib import pyplot as plt # we will also use the matplotlib
import warnings
warnings.simplefilter('ignore', category=UserWarning) # suppresses warni

# command below ensures matplotlib output can be included in Notebook

%matplotlib inline
```

```
In [2]: import os

# Get the current working directory
current_directory = os.getcwd()
print("Current working directory:", current_directory)
```

```
# Get the absolute path of "Training Data.csv"
file_path = os.path.abspath("Training Data.csv")
print("File absolute path:", file_path)
```

Current working directory: /Users/linlihsuan/Desktop/Linda/QMUL/Data Analysis/coursework1

File absolute path: /Users/linlihsuan/Desktop/Linda/QMUL/Data Analysis/coursework1/Training Data.csv

In [3]: `df = pd.read_csv('/Users/linlihsuan/Desktop/Linda/QMUL/Data Analysis/Lab_`

2.2 Exploring size of data and variable types

In [4]: `df.columns` # Listing the dataframe columns

Out[4]: Index(['Id', 'Income', 'Age', 'Experience', 'Married/Single',
'House_Ownership', 'Car_Ownership', 'Profession', 'CITY', 'STATE',
'CURRENT_JOB_YRS', 'CURRENT_HOUSE_YRS', 'Risk_Flag'],
dtype='object')

In [5]: `df.dtypes` # returns the datatype of each column – Pandas sets type from v

Out[5]:

Id	int64
Income	int64
Age	int64
Experience	int64
Married/Single	object
House_Ownership	object
Car_Ownership	object
Profession	object
CITY	object
STATE	object
CURRENT_JOB_YRS	int64
CURRENT_HOUSE_YRS	int64
Risk_Flag	int64
dtype:	object

In [6]: `df.shape` # returns the shape of data – 614 rows, and 13 columns

Out[6]: (252000, 13)

2.3 Preliminary exploration of values in the data

In [7]: `df.head(5)` # Let's view the first few rows of the dataframe

Out [7]:

	Id	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownership
0	1	1303834	23	3	single	rented	n
1	2	7574516	40	10	single	rented	n
2	3	3991815	66	4	married	rented	n
3	4	6256451	41	2	single	rented	ye
4	5	5768871	47	11	single	rented	n

In [8]: `df.describe()` # Let's view the description of the numerical values in the

Out [8]:

	Id	Income	Age	Experience	CURRENT_J
count	252000.000000	2.520000e+05	252000.000000	252000.000000	252000
mean	126000.500000	4.997117e+06	49.954071	10.084437	6
std	72746.278255	2.878311e+06	17.063855	6.002590	3
min	1.000000	1.031000e+04	21.000000	0.000000	0
25%	63000.750000	2.503015e+06	35.000000	5.000000	3
50%	126000.500000	5.000694e+06	50.000000	10.000000	6
75%	189000.250000	7.477502e+06	65.000000	15.000000	9
max	252000.000000	9.999938e+06	79.000000	20.000000	14

In [9]: `df.info()` # This pandas function returns the data types associated with e

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 252000 entries, 0 to 251999
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    252000 non-null int64
1   Income               252000 non-null int64
2   Age                 252000 non-null int64
3   Experience           252000 non-null int64
4   Married/Single      252000 non-null object
5   House_Ownership    252000 non-null object
6   Car_Ownership      252000 non-null object
7   Profession          252000 non-null object
8   CITY                252000 non-null object
9   STATE              252000 non-null object
10  CURRENT_JOB_YRS     252000 non-null int64
11  CURRENT_HOUSE_YRS  252000 non-null int64
12  Risk_Flag          252000 non-null int64
dtypes: int64(7), object(6)
memory usage: 25.0+ MB
```

2.4 Distributions of categorical values

```
In [10]: print('ID is non-numeric and is unique for each row: {} different values
          .format(len(df['Id'].unique()))
df = df.drop(['Id'], axis=1)
```

ID is non-numeric and is unique for each row: 252000 different values ...
so we can drop it

```
In [11]: print('Values and counts for Risk Flag are:\n{}'.format(df['Risk_Flag'].v
          print('Values and normalised counts for Risk_Flag are:\n{}'.format(df['Ri
```

Values and counts for Risk Flag are:

```
Risk_Flag
0    221004
1     30996
```

Name: count, dtype: int64

Values and normalised counts for Risk_Flag are:


```
Risk_Flag
0    0.877
1    0.123
```

Name: proportion, dtype: float64

2.5 Distributions of continuous variables


```
In [12]: fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(16, 4))
df['Income'].plot(kind='hist', bins=50, ax=axes[0], xlabel="Applicant Inc
          title="Histogram")
df['Income'].plot(kind='density', color='r', ax=axes[1], title='Density P
df['Income'].plot(kind='box', ax=axes[2], ylabel='Applicant Income ($)',
          xlabel='', title='Boxplot')

plt.show()
```




```
In [13]: fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(16, 4))
df['Age'].plot(kind='hist', bins=50, ax=axes[0], xlabel="Age",
               title="Histogram")
df['Age'].plot(kind='density', color='r', ax=axes[1], title='Density Plot')
df['Age'].plot(kind='box', ax=axes[2], ylabel='Age',
               xlabel='', title='Boxplot')

plt.show()
```



```
In [14]: fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(16, 4))
df['Age'].plot(kind='box', ax=axes[1], ylabel='Age',
               xlabel='', title='Boxplot')

df['Income'].plot(kind='box', ax=axes[0], ylabel='Income (M)',
                  xlabel='', title='Boxplot')
plt.show()
```



2.6 Bivariate Analysis - does a categorical variable look predictive of risk flag

```
In [15]: Married_Single=pd.crosstab(df['Married/Single'],df['Risk_Flag'])
print('Counts for Married/Single vs Loan_Status are:\n{}'.format(Married_

Married_Single = Married_Single.div(Married_Single.sum(1).astype(float),
print('\nNormalised counts for Married_Single vs Loan_Status are:\n{}'.fo

House_Ownership=pd.crosstab(df['House_Ownership'],df['Risk_Flag'])
print('Counts for House_Ownership vs Loan_Status are:\n{}'.format(House_0

House_Ownership = House_Ownership.div(House_Ownership.sum(1).astype(float)
print('\nNormalised counts for House_Ownership vs Loan_Status are:\n{}'.f

Car_Ownership=pd.crosstab(df['Car_Ownership'],df['Risk_Flag'])
print('Counts for Car_Ownership vs Loan_Status are:\n{}'.format(Car_Owner

Car_Ownership = Car_Ownership.div(Car_Ownership.sum(1).astype(float), axi
print('\nNormalised counts for Car_Ownership vs Loan_Status are:\n{}'.for

fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(14, 4))

Married_Single.plot(kind="bar",stacked=True, ax=axes[0])
House_Ownership.plot(kind="bar",stacked=True, ax=axes[1])
Car_Ownership.plot(kind="bar",stacked=True, ax=axes[2])
axes[0].tick_params(axis='x', rotation=45)
axes[1].tick_params(axis='x', rotation=45)
```

```
axes[2].tick_params(axis='x', rotation=45)
plt.show()
```

Counts for Married/Single vs Loan_Status are:

Risk_Flag	0	1
Married/Single		
married	23092	2636
single	197912	28360

Normalised counts for Married_Single vs Loan_Status are:

Risk_Flag	0	1
Married/Single		
married	0.897544	0.102456
single	0.874664	0.125336

Counts for House_Ownership vs Loan_Status are:

Risk_Flag	0	1
House_Ownership		
norent_noown	6469	715
owned	11758	1160
rented	202777	29121

Normalised counts for House_Ownership vs Loan_Status are:

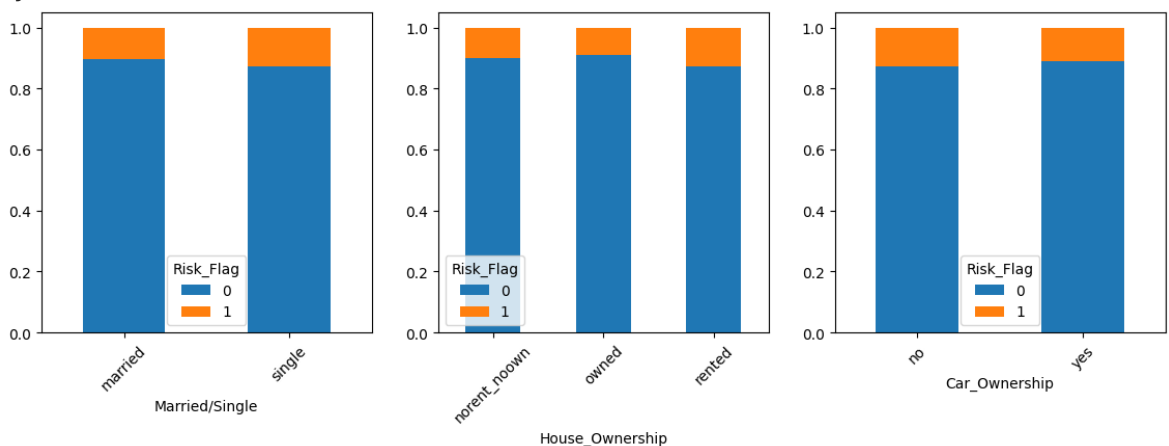
Risk_Flag	0	1
House_Ownership		
norent_noown	0.900473	0.099527
owned	0.910203	0.089797
rented	0.874423	0.125577

Counts for Car_Ownership vs Loan_Status are:

Risk_Flag	0	1
Car_Ownership		
no	153439	22561
yes	67565	8435

Normalised counts for Car_Ownership vs Loan_Status are:

Risk_Flag	0	1
Car_Ownership		
no	0.871812	0.128188
yes	0.889013	0.110987



```
In [16]: print(len(df.Profession.unique()))
print(len(df.STATE.unique()))
print(len(df.CITY.unique()))
```

51
29
317

3. Preprocessing data using Pandas

3.1 Missing value imputation

```
In [17]: df.isnull().sum()
```

```
Out[17]: Income                0
Age                0
Experience          0
Married/Single     0
House_Ownership    0
Car_Ownership      0
Profession         0
CITY               0
STATE             0
CURRENT_JOB_YRS    0
CURRENT_HOUSE_YRS  0
Risk_Flag          0
dtype: int64
```

```
In [18]: # No missing values need to be filled.
```

3.2 Convert categorical variables to integer ones

```
In [19]: print('Married/Single is originally a string variable:\n{}'.format(df['Ma
df['Married/Single'] = df['Married/Single'].astype('category').cat.codes
print('\nMarried/Single is converted to integer:\n{}'.format(df['Married/

# do the rest of the categorical variables
for n in ['House_Ownership', 'Profession', 'CITY', 'STATE', 'Car_Ownership']:

    df[n] = df[n].astype('category').cat.codes

print('\nData frame with categorical variables converted to integers:')
df.head()
```

Married/Single is originally a string variable:

Married/Single

single 226272

married 25728

Name: count, dtype: int64

Married/Single is converted to integer:

Married/Single

1 226272

0 25728

Name: count, dtype: int64

Data frame with categorical variables converted to integers:

```
Out[19]:
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownership	F
0	1303834	23	3	1	2	0	
1	7574516	40	10	1	2	0	
2	3991815	66	4	0	2	0	
3	6256451	41	2	1	2	1	
4	5768871	47	11	1	2	0	


4. Feature Selection

4.1 Look at correlations between each pair of variables using Pandas

Let's view the correlation between every pair of variables using Pandas

- used `***dataframe*.corr()*` - note only looks at numeric variables, so do this after you've converted categorical variables to integer

```
In [20]: corr=df.corr() # gives us the correlation values
plt.figure(figsize=(15,6))
sns.heatmap(corr, annot = True, cmap="BuPu") # let's visualise the corre
plt.xticks(rotation=45)
plt.show()
```



Observation

1. Consider dropping CURRENT_JOB_YRS as it is quite strongly correlated with Experience.

Let's drop some feature variables that seem to be highly correlated with other features and are therefore maybe not providing any useful extra information.


```
In [21]: # Let's drop the CURRENT_JOB_YRS features below
cols = ['CURRENT_JOB_YRS']
df = df.drop(columns=cols,axis=1)
```

4.2 Look at how strongly the remaining feature variables are associated with the target variable using Scikit-learn

This is an alternative/additional way of choosing some feature variables to drop using Scikit-Learn methods. The intuition is that feature variables which are not strongly associated with the target variable won't be very useful in predicting it and so are candidates to be dropped.

```
In [22]: # Split dataframe into feature variable inputs 'X' dataframe, and output
X = df.drop(['Risk_Flag'],axis=1)
y = df['Risk_Flag']
```

The first two methods below measure correlation between each feature variable and the target variable, but you can specify the correlation test. The chi2 (chi-squared) test is a standard linear correlation test, whereas the ftest is better at spotting non-linear correlations.

The **SelectKBest** method just returns the 'k' variables most highly correlated with the target, but here we are choosing to retrain them all.

```
In [23]: from sklearn.feature_selection import SelectKBest, chi2, mutual_info_classif
# let's call the k-best method with Chi-squared score and pass X and y as
chi2 = SelectKBest(score_func = chi2, k = 'all').fit(X,y)
# create Series with variable name as index, and scores as values, and so
chi2_sorted = pd.Series(data=chi2.scores_, index=X.columns).sort_values()
# Repeat but with other scoring functions
ftest = SelectKBest(score_func = f_classif, k = 'all').fit(X,y)
ftest_sorted = pd.Series(data=ftest.scores_, index=X.columns).sort_values()
mitest = SelectKBest(score_func = mutual_info_classif, k = 'all').fit(X,y)
mitest_sorted = pd.Series(data=mitest.scores_, index=X.columns).sort_valu
```

Instead of using correlation tests to show how strongly feature variables are associated with the target variable we can use some machine learning approaches which provide information about how strongly each target variable **as a by-product** of learning a predictive model. Here we use the ExtraTreesClassifier and MutualInformationClassifier.

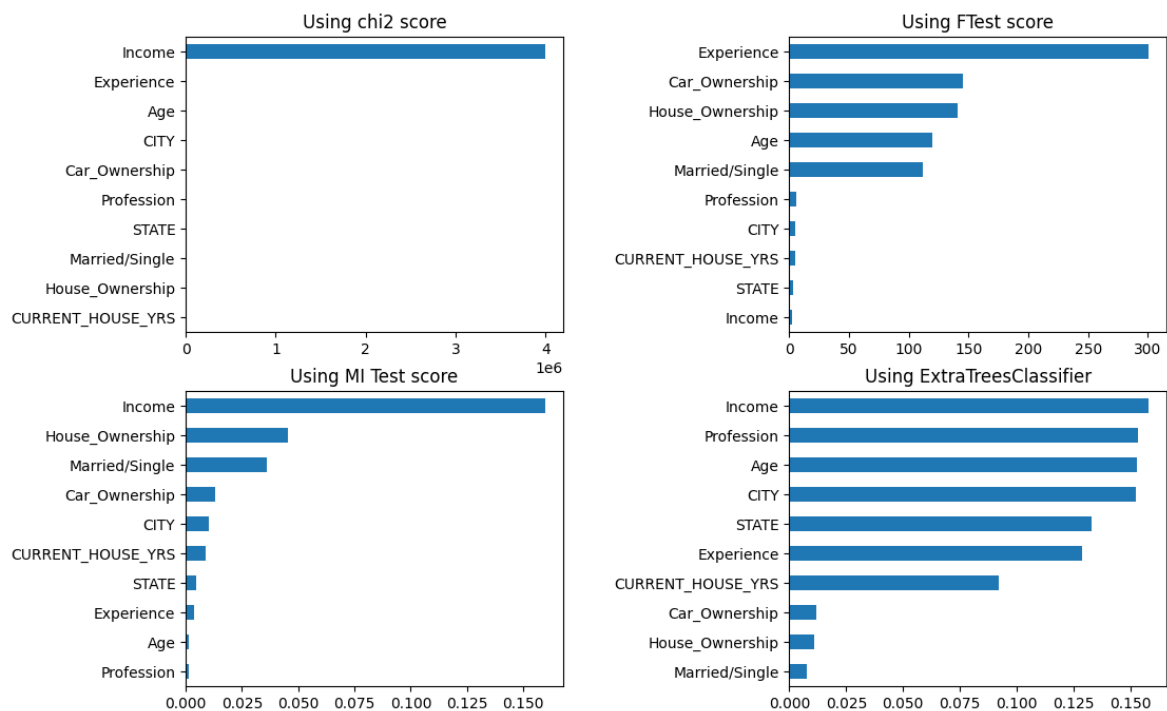
```
In [24]: from sklearn.ensemble import ExtraTreesClassifier # this is a method alte
xtrees = ExtraTreesClassifier().fit(X, y)
xtrees_sorted = pd.Series(data=xtrees.feature_importances_, index=X.colum
```

```
from sklearn.feature_selection import mutual_info_classif

muinfo = mutual_info_classif(X,y)
muinfo_sorted = pd.Series(data=muinfo, index=X.columns).sort_values()
```

Let's plot out the strength of association of the feature variables with the target variable using the four methods we've just run.

```
In [25]: fig, axes = plt.subplots(nrows=2, ncols=2, figsize=(12, 8))
plt.subplots_adjust(wspace=0.6)
chi2_sorted.plot(kind='barh', ax=axes[0, 0], title='Using chi2 score')
ftest_sorted.plot(kind='barh', ax=axes[0, 1], title='Using FTest score')
xtrees_sorted.plot(kind='barh', ax=axes[1, 1], title='Using ExtraTreesClassifier')
mitest_sorted.plot(kind='barh', ax=axes[1, 0], title='Using MI Test score')
plt.show()
```



Observation

1. Income gets a high score on all measures and so is very predictive feature.

Except for Current_Job_Yrs always being most strongly associated with Experience these four approaches don't have much else in common, so let's not delete any more feature variables.

5. Prediction and Evaluation using Scikit-Learn

```
In [26]: # for this, we will import another library so that we don't have to code
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, f1_score, precision_score, \
    recall_score, confusion_matrix
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

```

In [27]: def train_and_evaluate(model, X, y):
        """
        Train and evaluate a classification model on training data
        and produce accuracy metrics for a separate test set.
        """

        print('\nResults from algorithm {}'.format(model))

        # Split data into train and test – we will use test for the final acc
        # and not use it to train the model. This is good practice, particula
        # using cross-validation to select model parameters ... that way, the
        # of the test data don't leak into the model training

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0

        # Cross-validation accuracy gives an indication of variation in accur
        # estimate for overall accuracy than just a single estimate. The mean
        # accuracy is therefore a better guide when selecting model parameter
        # cross_val_score
        # scoring
        # cv=5
        scores = cross_val_score(model, X_train, y_train, cv=5, scoring='accu
        print('Mean cross-validation accuracy is {:.3f} with SD {:.3f}'
              .format(np.mean(scores), np.std(scores)))

        # Fit model using all of the reserved training data ... look at train
        # which we generally expect to be better than test accuracy

        learnt_model = model.fit(X_train, y_train)
        print('\nAccuracy on training data is {:.3f}\n'.format(model.score(X_

        # User predict() to predict target values from test feature variables
        # use functions to compute evaluation metrics relevant to binary outc

        y_pred = model.predict(X_test)
        print('Test data metrics: accuracy={:.3f}, f1={:.3f}, precision={:.3f}
              .format(accuracy_score(y_true=y_test, y_pred=y_pred),
                      f1_score(y_true=y_test, y_pred=y_pred),
                      precision_score(y_true=y_test, y_pred=y_pred),
                      recall_score(y_true=y_test, y_pred=y_pred)))

        # Draw out a confusion matrix

        cm = confusion_matrix(y_true=y_test, y_pred=y_pred)
        plt.figure(figsize=(3, 3))
        ax = sns.heatmap(cm,annot=True, xticklabels=['N', 'Y'], cbar=False,
                        yticklabels=['N', 'Y'], square=True,
                        linewidths=8.0, fmt='d') # plots the confusion matri
        ax.set_xlabel('Predicted Red Flag')
        ax.set_ylabel('Actual Red Flag')
        plt.show()

        return learnt_model

```

```

In [28]: from sklearn.model_selection import train_test_split
        from sklearn.model_selection import cross_val_score

```

```

In [29]: # and now try K-nearest neighbour

```

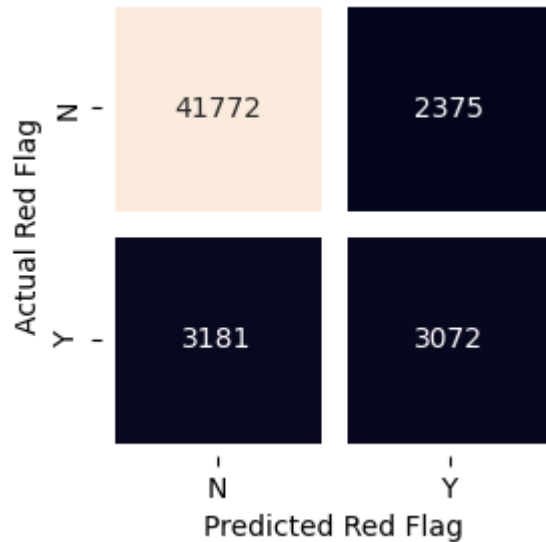
```
from sklearn.neighbors import KNeighborsClassifier
_ = train_and_evaluate(KNeighborsClassifier(), X, y)
```

Results from algorithm KNeighborsClassifier():

Mean cross-validation accuracy is 0.889 with SD 0.001

Accuracy on training data is 0.899

Test data metrics: accuracy=0.890, f1=0.525, precision=0.564, recall=0.491



In [31]: *# and finally a decision tree*

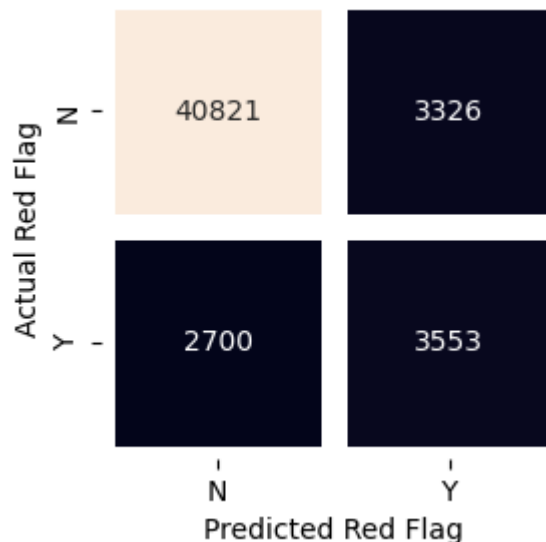
```
from sklearn.tree import DecisionTreeClassifier, plot_tree
learnt_model = train_and_evaluate(DecisionTreeClassifier(), X, y)
plt.figure(figsize=(16, 8))
plot_tree(learnt_model, max_depth=3, feature_names=list(X.columns),
          fontsize=9, filled=True, impurity=False, rounded=True)
plt.show()
```

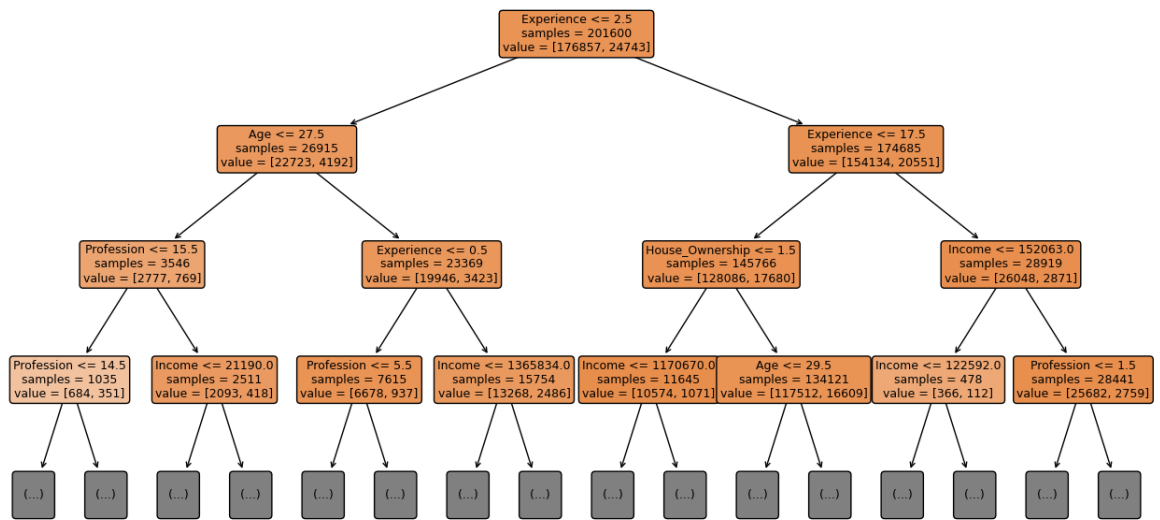
Results from algorithm DecisionTreeClassifier():

Mean cross-validation accuracy is 0.881 with SD 0.001

Accuracy on training data is 0.936

Test data metrics: accuracy=0.880, f1=0.541, precision=0.516, recall=0.568





In []: